

Christopher Enriquez

Data Engineering • Cloud Infrastructure • ETL Pipelines

+1 415-769-1152 | chris.fog.city@gmail.com | linkedin.com/in/chris-a-enriquez/ | github.com/BubbaLubz

EDUCATION

California Polytechnic State University

San Luis Obispo, CA

B.S. in Interdisciplinary Studies, Minors in Computer Science and Music

Expected June 2028

- **Relevant Coursework:** Data Structures, Discrete Math, Databases, Systems Programming
- **Leadership:** Sub-Section Leader, Cal Poly Drumline — led full percussion ensemble rehearsals, coached members, and maintained performance standards

EXPERIENCE

Data Engineering Lead

January 2026 - June 2026

EKI Environment and Water Inc.

San Luis Obispo, CA

- Led and coordinated a 5-person team building a Dockerized PFAS data standardization tool (OpenAI/Anthropic APIs) to normalize inconsistent federal/state schemas across 1M+ records. Improved non-detect identification from 0% to 81%, reduced QA false positives by 53%.
- Built Python ETL pipelines to clean, deduplicate, enrich, and validate 1M+ PFAS records, including batched geocoding for 10,000+ coordinates.
- Designed a normalized MySQL data model and SQL reporting layer powering Power BI, Tableau, and Streamlit dashboards for California PFAS trend analysis.

PROJECTS

Cloud Music Analytics Platform | *Python, AWS Lambda, Kafka, Athena, dbt, Grafana, Terraform*

- Built an event-driven AWS analytics pipeline ingesting 165k+ ListenBrainz events through Kafka, with 30-minute polling intervals and at-least-once delivery guaranteed via offset-after-write commits.
- Implemented checkpoint-based + Redis deduplication across 30-minute poll cycles to prevent reprocessing previously ingested listens.
- Provisioned infrastructure with Terraform including a 1 GB per-query Athena cost guardrail and 90-day S3 lifecycle policy on raw data.

Data Infrastructure Diagram Tool | *TypeScript, React, ReactFlow, Dagre, Anthropic API, Supabase* | **DEMO**

- Built an AI infrastructure diagramming tool that converts Claude-generated JSON into Dagre-computed ReactFlow graphs, using a 37-component system across 11 infrastructure categories.
- Reduced edit-mode payload size by 40% and input tokens by 34% by replacing full ReactFlow object serialization with minimal node/edge schemas.
- Detected ambiguous infrastructure prompts with 100% accuracy on a 20-prompt test set, with a user-triggered best-effort fallback to avoid clarification loops.

StickSplit (KiroHacks 26) | *Python, FastAPI, MediaPipe, OpenCV, Firebase Firestore, Docker* | **DEMO**

- Built a video data pipeline that ingests drumming videos, samples frames with OpenCV, extracts 19 MediaPipe landmarks, and computes 6 biomechanical metrics.
- Calibrated landmark thresholds across 10 test videos, reducing pose-signal saturation from 80% to 30% and correcting metric inflation by 85%.
- Developed a Dockerized FastAPI microservice for video ingestion, landmark extraction, metric computation, Firestore writes, and Pydantic validation.

TECHNICAL SKILLS

Languages: Python, SQL, TypeScript

Data Engineering: ETL/ELT Pipelines, Data Modeling, Schema Design, Streaming Pipelines, REST API Integration

Cloud & Infra: AWS Lambda, S3, EventBridge, CloudWatch, Terraform, Upstash Kafka, Redshift Serverless

Databases: PostgreSQL, MySQL, SQLite, DuckDB, Firebase Firestore

Orchestration & Quality: dbt Core, Prefect Cloud, Great Expectations

Libraries: Pandas, NumPy, OpenCV, Pydantic, SQLAlchemy, FastAPI, MediaPipe, ReactFlow, Dagre

AI/APIs: Anthropic API, OpenAI API, Spotify API

Tools: Docker, Git, Linux, Power BI, Tableau, Grafana